

Reproducible Documents with { rmarkdown }

Day 1

Instructor: Selina Baldauf

Freie Universität Berlin - Theoretical Ecology


2022-22-03 (updated: 2022-03-31)



Topics today

- Introduction to the package
- Basic workflow
- First R Markdown document

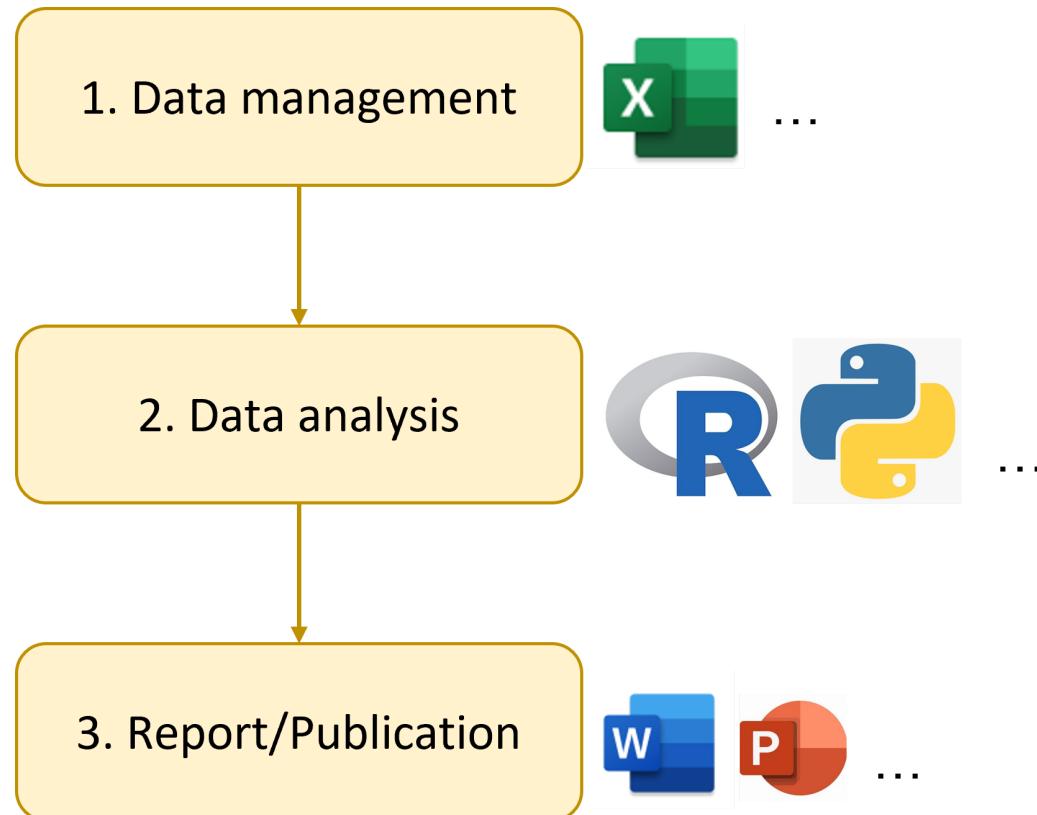
A standard workflow



Hard to answer questions:

- How did you produce this figure? What analysis is behind it?
- Where does this value come from?
- I found an error in the raw data. Can you repeat the analysis?

A standard workflow



Main problem with this workflow: **error prone** and **non-reproducible** (and it can be annoying)

If you have to repeat the analysis

- Redo all figures and tables
- Update document and presentation manually
 - Manual copy pasting of values is very error prone
- You probably have to repeat this several times

Solution: A workflow using R Markdown

Basic idea: Have everything (code, text, metadata) in one place. Let `{rmarkdown}` do the magic:

- Run code and add output
- Return a nice document of the desired output format

Motivation for using R Markdown

- Reproducibility
 - Easy to redo analysis
 - Easy to verify and check
 - No more copy pasting
 - Continuous workflow that is independent of the person that wrote the workflow (no clicking involved)
- Documentation/Text, Code & Output in one place
- Use R pipeline to produce documents
 - Create an automatic workflow
- Fun

Example use cases

Any type of document, especially if it contains something produced by code

- Data analysis reports
 - For yourself
 - For your supervisors
 - ...
- Publication + Publication of analysis
 - E.g. also good for publishing code
- Presentations
- Websites
- Books
- ...

The R Markdown universe

- Many packages that provide additional functionality
 - Additional output formats
 - Templates
 - Formatting tools
 - Printing tools
 - ...

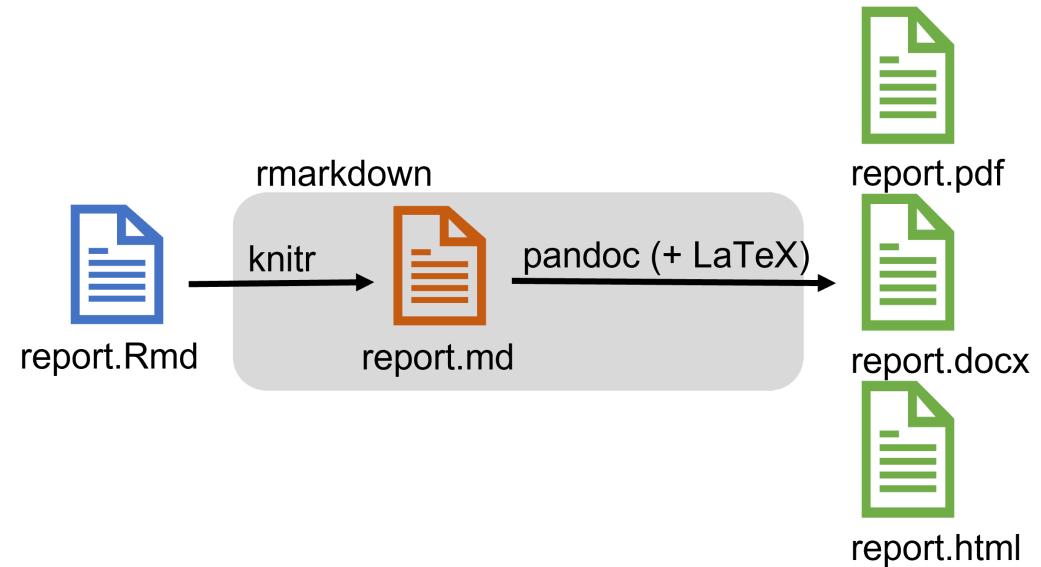
Just some examples:



There are many more ...

The basic workflow

1. Create an `.Rmd` document
2. Write text and R code into the document
3. Render the document to a defined output format using `rmarkdown`



Basic workflow in practice

Step 1: Create a new `.Rmd` document

- **Empty:** Create a new file in your project and save it with file ending `.Rmd`
- **From a template:**

R Markdown itself comes with some templates

- `File -> New File -> Rmarkdown...`

Additional packages come with additional packages

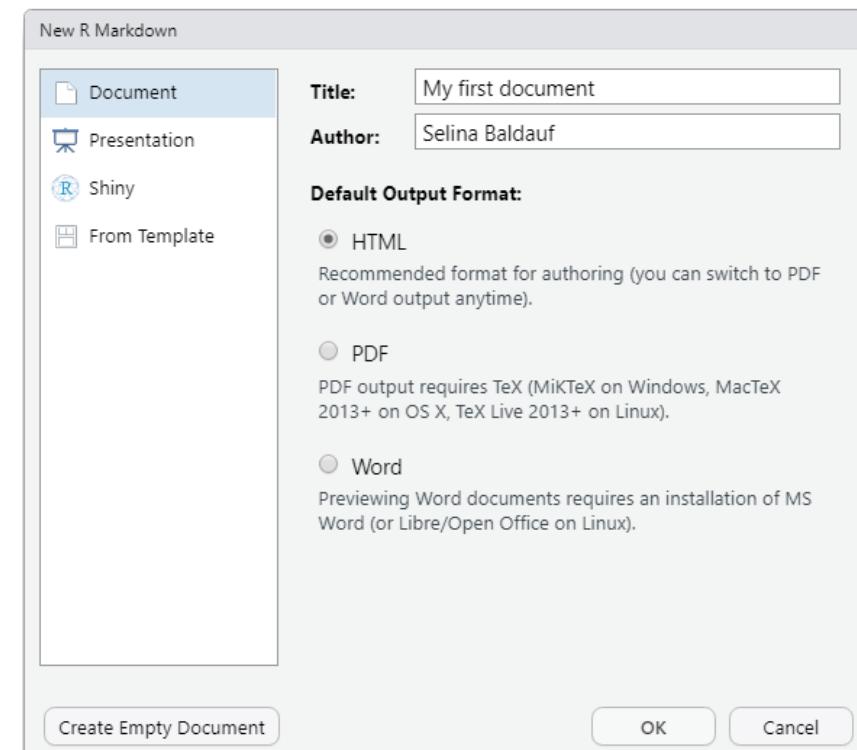
- `{rticles}` for article templates
- `{xaringan}` for presentations ...

Basic workflow in practice

Step 1: Create a new `.Rmd` document

File -> New File -> Rmarkdown . . .

- Select template on left
- Add title and author metadata
 - Can also be left blank and done later
- Select output format
- Click OK



Structure of an R Markdown document

```
1 ---  
2 title: "My first document"  
3 author: "Selina Baldauf"  
4 date: "3/22/2022"  
5 output: pdf_document  
6 ---  
7  
8 ```{r setup, include=FALSE}  
9 knitr::opts_chunk$set(echo = TRUE)  
10 ````  
11 ## R Markdown  
12  
13 This is an R Markdown document. Markdown is a simple formatting syntax for authoring  
HTML, PDF, and MS Word documents. For more details on using R Markdown see  
http://rmarkdown.rstudio.com.  
14  
15 When you click the **Knit** button a document will be generated that includes both  
content as well as the output of any embedded R code chunks within the document. You can  
embed an R code chunk like this:  
16  
17 ```{r cars}  
18 summary(cars)  
19 ````  
20  
21 ## Including Plots  
22  
23 You can also embed plots, for example:  
24  
25  
26 ```{r pressure, echo=FALSE}  
27 plot(pressure)  
28 ````  
29  
30 Note that the `echo = FALSE` parameter was added to the code chunk to prevent printing  
of the R code that generated the plot.  
31
```

Basic workflow in practice

Step 2: Write your document

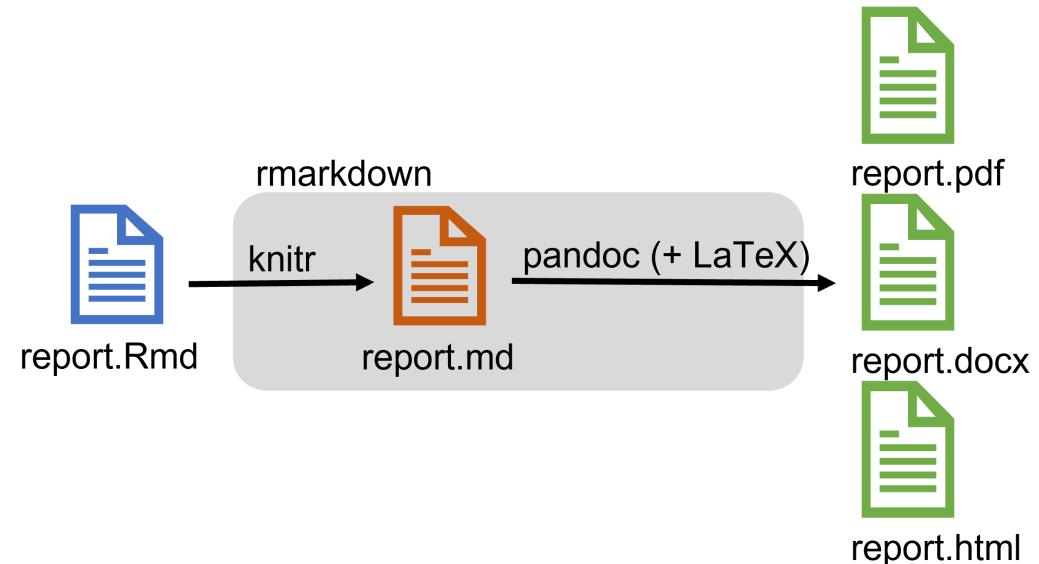
- Metadata
- Text
- R code

Basic workflow in practice

Step 3: Render/Knit the document to the desired output format

A multipstep process that is coordinated by the `{rmarkdown}` package

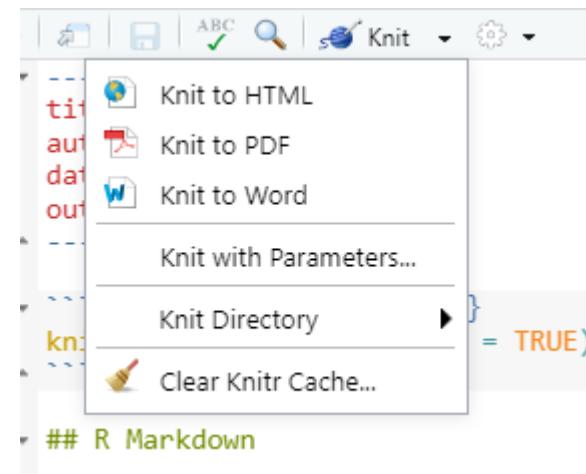
1. `{knitr}` (an R package) runs the code and adds the output to the text
 2. `knitr` creates a `.md` document
 3. `Pandoc` (+ LaTeX) converts the markdown document to the desired output format
- Document is knitted in a new R session
 - This makes sure that the document is reproducible and that it does not depend on the current environment



Basic workflow in practice

Step 3: Render/Knit the document to the desired output format

- Click the Knit Button
- Use the keyboard shortcut `Ctrl/Cmd + Shift + K`
- For more options, click the little arrow next to the Knit button



→ Iterate through steps 2 and 3 until finished

I recommend to knit often, otherwise it can become a pain to debug.

The text body: Markdown

The text body

- Text body in Markdown syntax
- Markdown is simple markup language to create formatted text
- Rmarkdown uses pandoc's markdown syntax
 - Find a full documentation [here](#)

The text body

The basics

- Bold: `**text**` becomes **text**
- Italic: `*text*` becomes *text*
- Subscript: `H~3~PO~4~` becomes H_3PO_4
- Superscript: `Cu^2+^` becomes Cu^{2+}

The text body

Code blocks

- Inline code: fenced off with 1 backtick becomes `code`
- Code blocks: fenced off with 3 backticks

```
```  
code
```
```

becomes

```
code
```

The text body

- Links: [text] (link)
 - [RStudio] (<https://www.rstudio.com>) becomes RStudio

The text body

- Include image (from file/web): ! [Figure caption] (some/image)
 - ! [Markdown logo]
(<https://upload.wikimedia.org/wikipedia/commons/thumb/4/48/Markdown-mark.svg/1200px-Markdown-mark.svg.png>) becomes



Markdown logo

- Footnote: ^ [a footnote]
 - Markdown logo by Dustin Curtis¹ [[available on Wikimedia](https://en.wikipedia.org/wiki/Markdown#/media/File:Markdown-mark.svg)]
(<https://en.wikipedia.org/wiki/Markdown#/media/File:Markdown-mark.svg>) becomes
 - Markdown logo by Dustin Curtis¹

¹ available on [Wikimedia](#)

The text body

Headers

```
# First level header  
## Second level header  
### Third level header
```

The text body

Itemized lists

- item 1
 - another item
- item 2
- item 3

becomes

- item 1
 - another item
- item 2
- item 3

Numbered lists

1. item 1
2. item 2
3. item 3

The text body

Math expressions can be written in LaTeX syntax

- Inline (enclosed with \$)

```
$f(k) = {n \choose k} p^k (1-p)^{n-k}$ returns  $f(k) = \binom{n}{k} p^k (1 - p)^{n-k}$ 
```

- As a separate block (enclosed with \$\$)

$$f(k) = \binom{n}{k} p^k (1 - p)^{n-k}$$

The text body

You don't need to remember all of this. [Here](#) is a quick reference for the most important things.

Always use spaces around markdown objects so that they can be rendered correctly, e.g.

```
## My section
```

This is an ordered list:

1. item 1
2. item 2

instead of

```
## My section
```

This is an ordered list:

1. item 1
2. item 2

The text body

Let's look at an example together

Can you identify the formatting elements in [this document](#)?

Now you

Task 1: Create a first docuemnt and add some text (45 min)

Find the task description [here](#)

The R code

The R code

R code can be included in **code chunks** as **inline code**

- R code can be displayed or hidden in output document
 - Inline code is hidden
- Code chunks can contain any type of R code
- R code is (by default) executed and output is included in document
 - Text output
 - Figures
 - Prepare future data analysis
 - ...

The R code

Code chunks starts and ends with 3 backticks

```
```{r}
code goes here
```
```

Example

```
```{r}
mean(1:3)
```
```

looks like this:

```
mean(1:3)

## [1] 2
```

The R code

Inline code starts and ends with 1 backtick

```
## `r `
```

Example

```
The mean of the values 1, 2 and 3 is `r mean(1:3)`
```

looks like this:

The mean of the values 1, 2 and 3 is 2.

The R code

Insert a code chunk

- Insert a code chunk by going to `Code -> Insert chunk`
- Use the keyboard shortcut `Ctrl + Alt + I / Cmd + Option + I`
- Inline chunks have to be typed

The R code

Chunk names

- Code chunks can have names (but they don't have to)
- Names are added inside the code chunk

```
```{r calculate-mean}
mean(1:3)
```
```

- Easier to debug
- Code chunks appear in document outline

The R code

Run code chunk

- Code chunks are evaluated by `knitr` when rendering the document
- Code chunks can also be run like normal R code
- Run Code chunk by clicking on the green arrow next to the chunk

```
```{r cars}
summary(cars)
```
```

Chunk options

- Code chunks options give you fine control over the behavior of a chunk
- Chunk options are separated by commas

```
```{r calculate-mean, warning=TRUE, message=TRUE}
mean(1:3)
```
```

- Chunk options have a `name` and a `value`
 - Chunk options have default values

Chunk options

Some important general options

- `message`: TRUE, FALSE, Show message in output?
- `warning`: TRUE, FALSE, Show warning in output?
- `eval`: TRUE, FALSE, Evaluate the chunk?
- `echo`: TRUE, FALSE, Show source code in output?
- `include`: TRUE, FALSE, Include anything from the chunk in the output?
 - Runs the code but does not show any output or source code

Chunk options

Some important figure related options

- `fig.width` and `fig.height`: Size of graphical device in inches (i.e. size of the plots)
- `out.width` and `out.height`: Scale output of R plots, e.g. to scale images `out.width = "80%"`
- `fig.align`: Plot alignment, one of `"left"`, `"center"`, `"right"`
- `fig.cap`: A figure caption as a string

```
```{r important-figure, fig.cap="This is a nice plot.", fig.width=4, fig.height=4}
plot(1:10, 1:10)
```
```

Chunk options

- Some chunk options can be set by clicking on the gear icon next to the chunk



The screenshot shows an RStudio interface. On the left, there is a code editor window containing the following R code:

```
```{r cars, message=TRUE, warning=TRUE}
summary(cars)
````
```

On the right, a gear icon is visible, which opens a dropdown menu titled "Chunk options". This menu contains the following settings:

- Chunk Name: cars
- Output: (Use document default)
- Show warnings
- Show messages
- Use paged tables
- Use custom figure size

At the bottom of the menu are "Revert" and "Apply" buttons.

- See [here](#) for a comprehensive list of all chunk options

The setup chunk

- Set default chunk options in the beginning of the document
- Default options can be overwritten in individual chunks
- Default chunk options can be set via `knitr` in an R code chunk:

```
```{r setup, include=FALSE}
knitr::opts_chunk$set(
 echo = TRUE,
 warning = FALSE
)
```

```

Other language engines

You can also run code chunks of other languages:

```
names(knitr::knit_engines$get())
```

```
## [1] "awk"          "bash"         "coffee"        "gawk"         "groovy"       "haskell"      "lein"
"mysql"
## [9] "node"         "octave"       "perl"          "pgsql"        "Rscript"      "ruby"         "sas"
"scala"
## [17] "sed"          "sh"           "stata"         "zsh"          "asis"         "asy"          "block"
"block2"
## [25] "bslib"        "c"             "cat"           "cc"           "comment"     "css"          "dot"
"embed"
## [33] "fortran"      "fortran95"    "go"            "highlight"   "js"           "julia"        "python"
"R"
## [41] "Rcpp"         "sass"         "scss"          "sql"          "stan"         "targets"      "tikz"
"verbatim"
## [49] "theorem"      "lemma"        "corollary"    "proposition" "conjecture"  "definition"
"example"      "exercise"
## [57] "hypothesis"   "proof"        "remark"       "solution"
```

See [here](#) if you are interested

Now you

Task 2: Add some R code to the document (40 min)

Find the task description [here](#)

The YAML (/jæməl/) header

YAML (/ˈjæməl/) header

- YAML is a data-serialization language
- Goes between `---` on top
- Used for
 - Meta data
 - Document output formats and their options
- Is used by Pandoc, `rmarkdown` and `knitr`
- You don't have to have a YAML header for the document to work
 - But usually you have at least some meta data or options to specify

YAML (/ˈjæməl/) header

Meta data

```
---
```

```
title: "My first document"
subtitle: "Whatever subtitle makes sense"
author: "Selina Baldauf"
date: "`r Sys.Date()`"
```

```
---
```

- Inline R code can print the current date at knitting time

YAML (/ˈjæməl/) header

Document outputs

- html_document
- pdf_document
- word_document
- beamer_presentation
- powerpoint_presentation
- ...

To use the default version of an output format add

```
output: html_document
```

YAML (/ˈjæməl/) header

Multiple document outputs

- You can specify multiple document outputs in the YAML header.
- You have to list all output formats **inside** the output option
- Here just with default options for output types:

```
---
title: "My first document"
author: "Selina Baldauf"
date: "3/22/2022"
output:
  html_document: default
  pdf_document: default
  word_document: default
---
```

YAML (/ˈjæməl/) header

The YAML Syntax rules

```
---
```

```
title: "My first document"
author: "Selina Baldauf"
date: "3/22/2022"
output:
  html_document:
    fig_caption: true
  pdf_document: default
  word_document: default
---
```

- YAML Syntax: Indentation is important
 - Two spaces or a tab for lists
 - Values follow a :
- Order of options does not matter
- Translation R - YAML: TRUE, FALSE, NULL become true/yes, false/no, null
- Strings can be quoted but don't have to (if they don't contain special characters)

YAML (/ˈjæməl/) header

Document output options

- Every output format has options than can be specified in the YAML header
 - Some options are shared between formats
 - Some options are specific to formats
- Have a look at `?rmarkdown::pdf_document`, `?rmarkdown::html_document`, etc. to see all options and their default values
- Use the [R Markdown cheat sheet](#) as reference for the most important options

YAML (/ˈjæməl/) header

Document output options for all 3 output types

```
output:  
  html_document:  
    toc: true  
    toc_depth: 2  
    number_sections: true  
    fig_caption: true  
    highlight: "espresso"  
    df_print: "kable"
```

- `toc`: Show table of contents?
- `toc_depth`: Lowest header level for table of contents
- `number_sections`: Should sections be numbered?
- `fig_caption`: Render figure captions?
 - For pdf this option also numbers the figures
- `highlight`: Which style to use for syntax highlighting?
 - Some options: `"kate"`, `"tango"`, `"zenburn"`
- `df_print`: How should tables be printed?
 - `"kable"` is nicer than the default printing of tables

YAML (/ˈjæməl/) header

Document output options for all 3 output types

```
output:  
  html_document:  
    toc: true  
    toc_depth: 2  
    number_sections: true  
    fig_caption: true  
    highlight: "espresso"  
    df_print: "kable"  
  word_document:  
    toc: true  
    toc_depth: 2  
    number_sections: true  
    fig_caption: true  
    highlight: "kate"  
    df_print: "kable"  
  pdf_document:  
    toc: true  
    fig_caption: true
```

- Decide which output `rmarkdown` should render
 - Choose the output type with the little arrow next to the knit button
- If you knit without specifying the output type, the last rendered type is taken
- `knitr` will look at the options for the document type you chose before rendering

YAML (/ˈjæməl/) header

Document output options for `html_document` only

```
output:  
  html_document:  
    toc: true  
    toc_float: true  
    code_folding: "hide"  
    code_download: true  
    theme: "flatly"
```

- `toc_float`: Float the table of content on the side?
- `code_folding`: e.g. "hide" to hide code and make it available via a button
 - interesting in combination with `echo=TRUE` as default chunk option
- `code_download`: Integrate a download button for the code?
- `theme`: Apply a theme to the output, e.g. "flatly", "journal" or "united"
 - [Here](#) you can find a list of other themes to try

YAML (/ˈjæməl/) header

Document output options for `pdf_document` only

```
---
```

```
output:
  pdf_document:
    fig_caption: true
    citation_package: "natbib"
    keep_tex: true
    template: "my_template.tex"
```

```
---
```

- `fig_caption` will automatically number figures
- `citation_package` define the citation package to use (e.g. `"natbib"` or `"biblatex"`)
- `keep_tex`: Keep the intermediate `*.tex` file?
- `template`: Path to a template file

YAML (/ˈjæməl/) header

Document output options for `pdf_document` only

```
---
```

```
output:
  pdf_document:
    fig_caption: true
    citation_package: "natbib"
    keep_tex: true
    template: "my_template.tex"
  fontsize: 11pt
  geometry: "margin=1in"
  documentclass: "article"
  urlcolor: "blue"
---
```

- Define some LaTeX options as *top-level* YAML metadata
- These options will be ignored when knitting to a different format

YAML (/ˈjæməl/) header

Document output options for `word_document` only

```
---
```

```
title: "My first document"
author: "Selina Baldauf"
date: "3/22/2022"
output:
  word_document:
    reference_docx: "my-styles.docx"
---
```

- Not so many features directly available in R Markdown
- `reference_docx`: Use an office template for customization
 - [Read](#) or [watch](#) how to create a custom office template
- Package `{officedown}` might be useful
 - It offers templates for advanced Word documents and Powerpoint presentations
 - `File -> New File -> R Markdown... -> From Template`

Now you

Task 3: Change meta data and document options (30 min)

Find the task description [here](#)