

Reproducible Documents with { rmarkdown }

Day 1

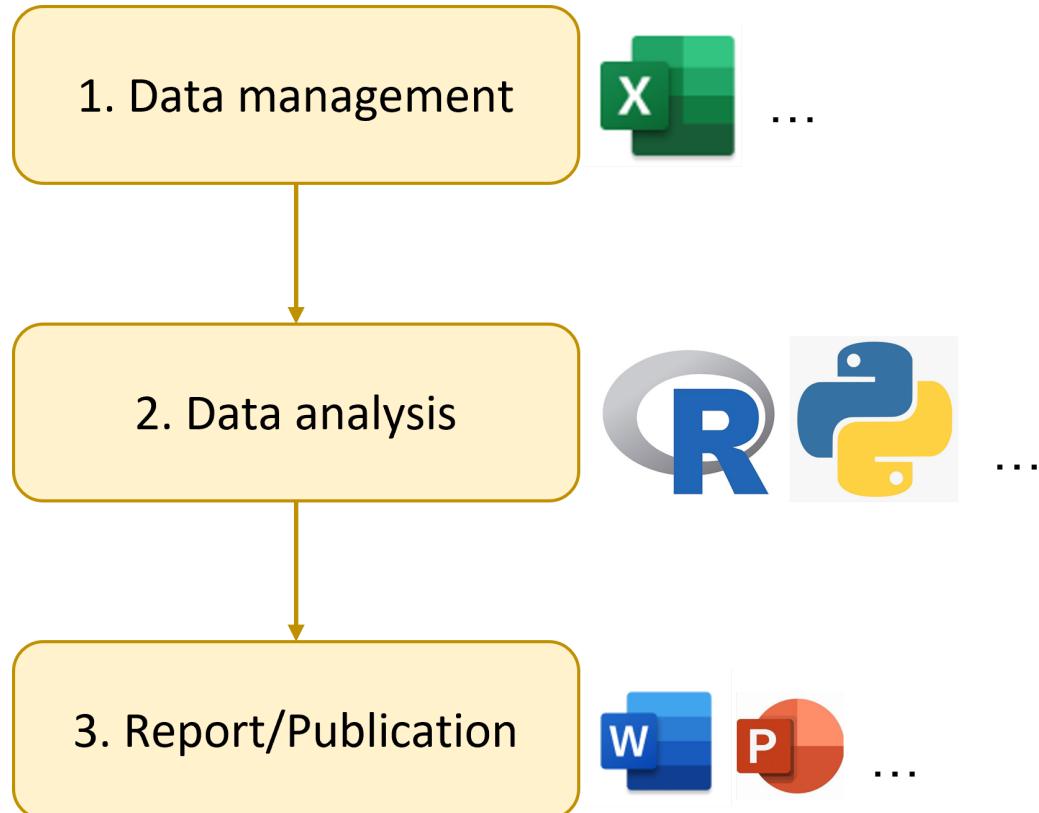
Instructor: Selina Baldauf

Freie Universität Berlin - Theoretical Ecology

2022-22-03 (updated: 2022-03-28)



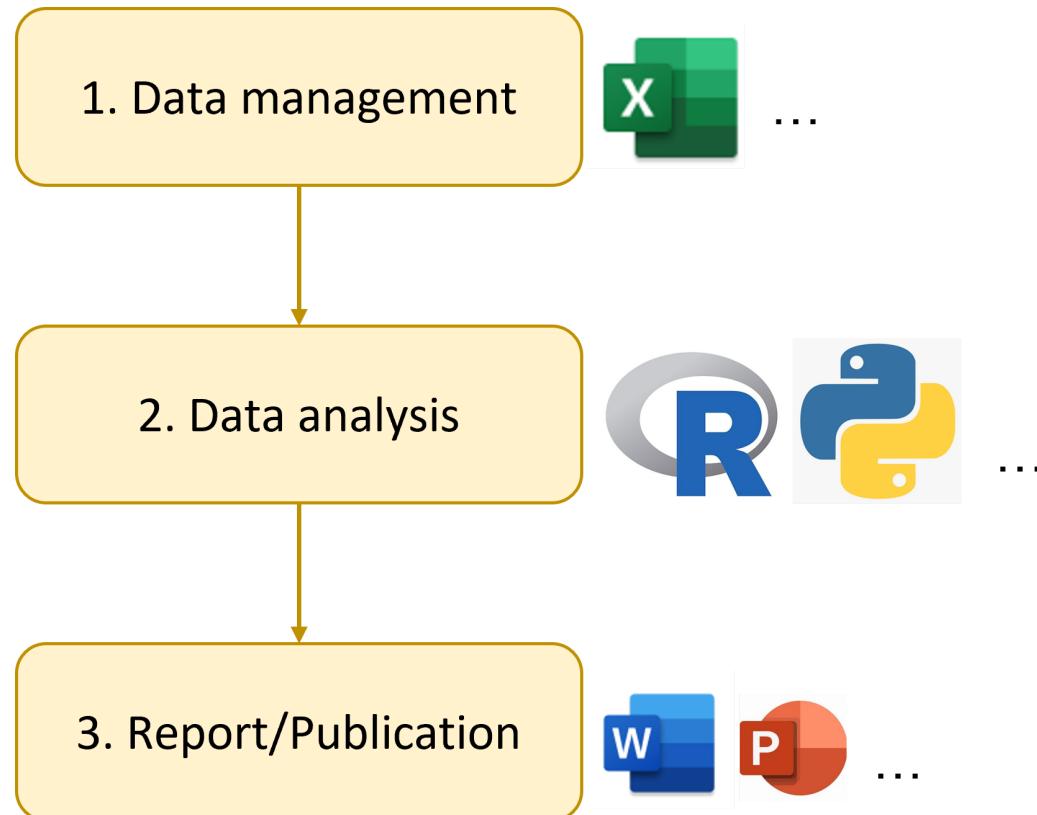
A standard workflow



Hard to answer questions:

- How did you produce this figure? What analysis is behind it?
- What data did you use for this table? Did you leave any data points out?
- Where does this value come from?
- I found an error in the raw data. Can you repeat the analysis?

A standard workflow



Main problem with this workflow: **error prone** and **non-reproducible**

If you have to repeat the analysis

- Redo all figures and tables
- Update document and presentation manually
 - Manual copy pasting of values is very error prone
- Annoying and you probably have to repeat this several times

Solution: A workflow using `{rmarkdown}`

Basic idea: Have everything (code, text, metadata) in one place. Let `{rmarkdown}` do the magic:

- Run code and add output to the document
- Return a nice document of the desired output format

Motivation for using {rmarkdown}

- Reproducibility
 - Easy to redo analysis
 - Easy to verify and check
 - No more copy pasting
 - Continuous workflow that is independent of the person that wrote the workflow (no clicking involved)
- Documentation/Text, Code & Output in one place
- Use R pipeline to produce documents
 - Create an automatic workflow
- Fun

Example use cases

Any type of document, especially if it contains something produced by code

- Data analysis reports
 - For yourself
 - For your supervisors
 - ...
- Publication + Publication of analysis
 - E.g. also good for publishing code
- Presentations
- Websites
- Books
- ...

The `rmarkdown` universe

- Many packages that provide additional functionality
 - Additional output formats
 - Templates
 - Formatting tools
 - Printing tools
 - ...

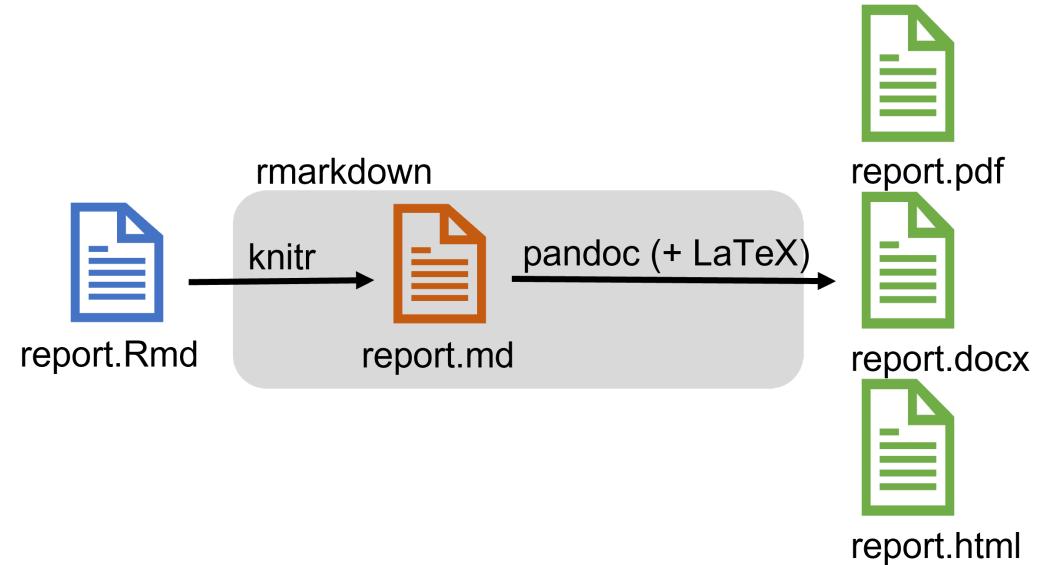
Just some examples:



There are many more ...

The basic workflow

1. Create an `.Rmd` document
2. Write text and R code in the document
3. Render the document to a defined output format using `rmarkdown`



Basic workflow in practice

Step 1: Create a new `.Rmd` document

- **Empty:** Create a new file in your project and save it with file ending `.Rmd`
- **From a template:**

Rmarkdown itself comes with some templates

- `File -> New File -> Rmarkdown...`

Additional packages come with additional packages

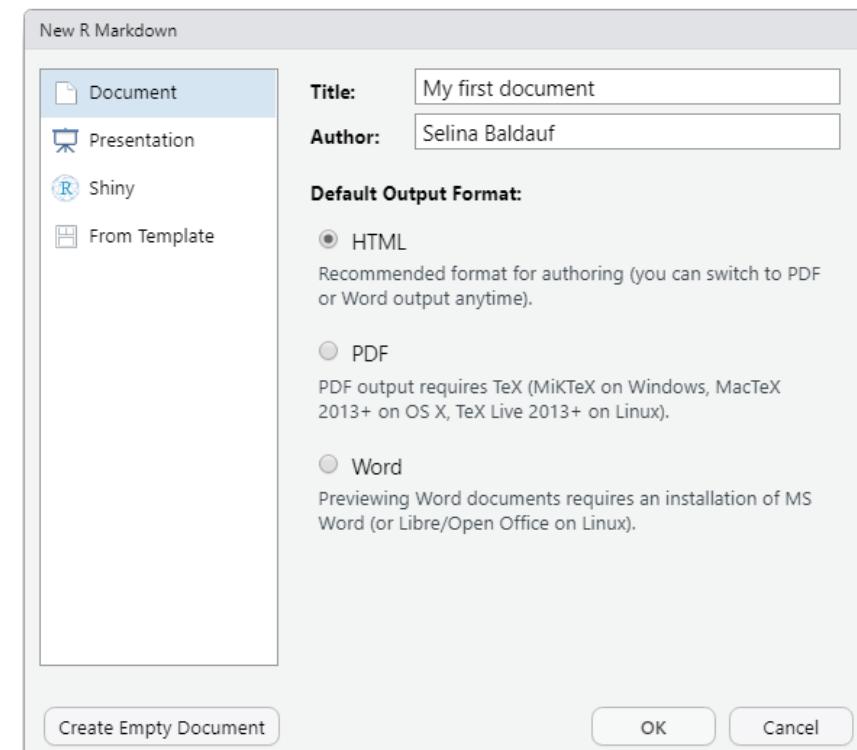
- `{rticles}` for article templates
- `{xaringan}` for presentations ...

Basic workflow in practice

Step 1: Create a new `.Rmd` document

File -> New File -> Rmarkdown . . .

- Select template on left
- Add title and author metadata
 - Can also be left blank and done later
- Select output format
- Click OK



Structure of an `rmarkdown` document

The screenshot shows the RStudio interface with an R Markdown file open. The code editor contains the following content:

```
1 ---  
2 title: "My first document"  
3 author: "Selina Baldauf"  
4 date: "3/22/2022"  
5 output: pdf_document  
6 ---  
7  
8 ```{r setup, include=FALSE}  
9 knitr::opts_chunk$set(echo = TRUE)  
10 ````  
11 ## R Markdown  
12  
13 This is an R Markdown document. Markdown is a simple formatting syntax for authoring  
14 HTML, PDF, and MS Word documents. For more details on using R Markdown see  
15 <http://rmarkdown.rstudio.com>.  
16 When you click the **Knit** button a document will be generated that includes both  
17 content as well as the output of any embedded R code chunks within the document. You can  
18 embed an R code chunk like this:  
19  
20 ```{r cars}  
21 summary(cars)  
22 ````  
23 ## Including Plots  
24  
25 You can also embed plots, for example:  
26  
27 ```{r pressure, echo=FALSE}  
28 plot(pressure)  
29 ````  
30 Note that the `echo = FALSE` parameter was added to the code chunk to prevent printing  
31 of the R code that generated the plot.
```

Annotations on the left side of the code editor highlight specific sections:

- YAML header with metadata**: Points to the first six lines of the file, which define the document's title, author, date, and output type.
- R code in code chunks**: Points to the R code chunks enclosed in triple backticks, specifically lines 8 through 10 and line 19.
- Text in Markdown syntax**: Points to the explanatory text and code examples written in Markdown syntax, starting from line 14.

Basic workflow in practice

Step 2: Write your document

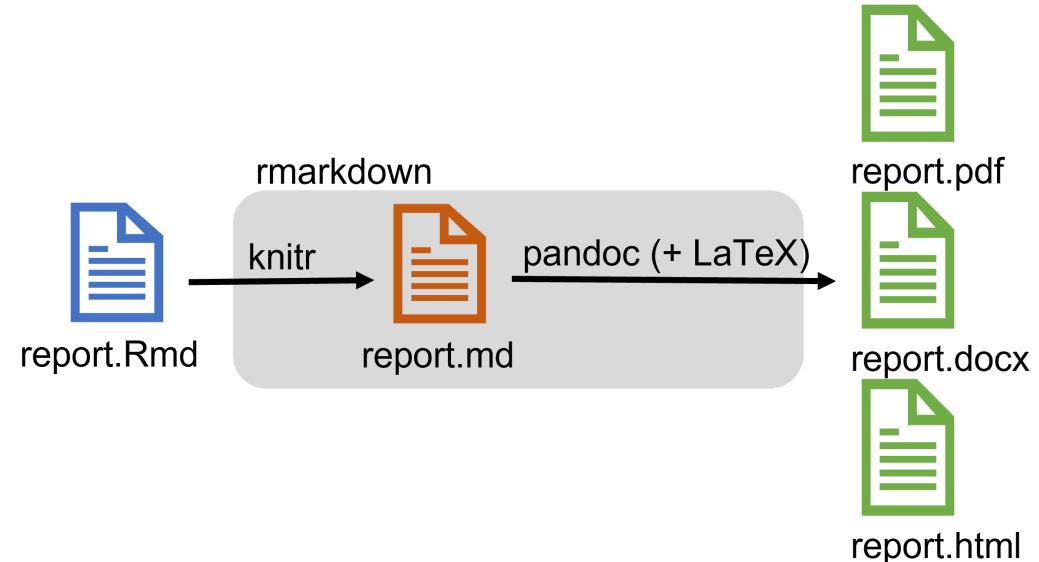
- Metadata
- Text
- R code

Basic workflow in practice

Step 3: Render/Knit the document to the desired output format

A multipstep process that is coordinated by the `rmarkdown` package

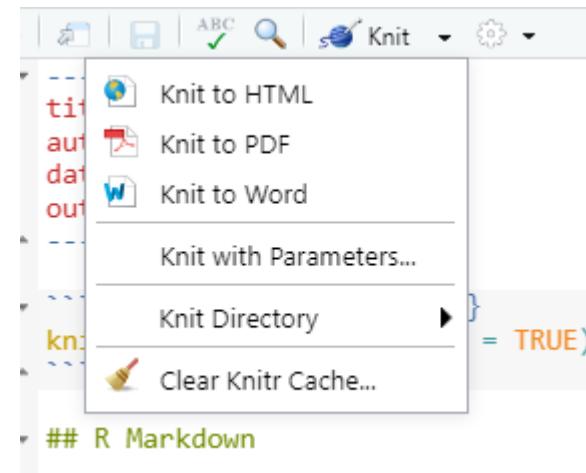
- `{knitr}` (an R package) runs the code and adds the output to the text
- `knitr` creates a `.md` document
- `Pandoc` (+ LaTeX) converts the markdown document to the desired output format
- Document is knitted in a new R session
 - This makes sure that the document is reproducible and that it does not depend on the current environment



Basic workflow in practice

Step 3: Render/Knit the document to the desired output format

- Click the Knit Button
- Use the keyboard shortcut `Ctrl/Cmd + Shift + K`
- For more options, click the little arrow next to the Knit button



→ Iterate through steps 2 and 3 until finished

I recommend to knit often, otherwise it can become a pain to debug.

The text body

The text body

- Text body in Markdown syntax
- Markdown is simple markup language to create formatted text
- Rmarkdown uses pandoc's markdown syntax
 - Find a full documentation [here](#)

The text body

The basics

- Bold: `**text**` becomes **text**
- Italic: `*text*` becomes *text*
- Subscript: `H~3~PO~4~` becomes H_3PO_4
- Superscript: `Cu^2+^` becomes Cu^{2+}

The text body

Code blocks

- Inline code: becomes `code`
- Code blocks: fenced off with 3 backticks

```
```  
code
```
```

becomes

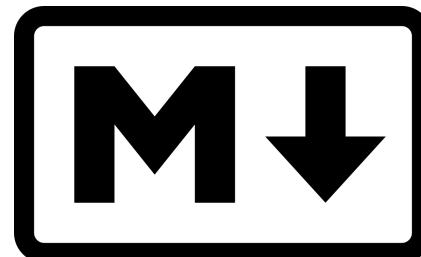
```
code
```

The text body

- Links: [text] (link)
 - [RStudio] (<https://www.rstudio.com>) becomes RStudio

The text body

- Include image (from file/web): ! [Alt text/Figure caption] (some/image)
 - ! [Markdown logo]
(<https://upload.wikimedia.org/wikipedia/commons/thumb/4/48/Markdown-mark.svg/1200px-Markdown-mark.svg.png>) becomes



Markdown logo

- Footnote: ^ [a footnote]
 - Markdown logo by Dustin Curtis¹ [[available on Wikimedia](https://en.wikipedia.org/wiki/Markdown#/media/File:Markdown-mark.svg)]
(<https://en.wikipedia.org/wiki/Markdown#/media/File:Markdown-mark.svg>) becomes
 - Markdown logo by Dustin Curtis¹

¹ available on [Wikimedia](#)

The text body

Headers

```
# First level header  
## Second level header  
### Third level header
```

The text body

Itemized lists

- item 1
 - another item
- item 2
- item 3

becomes

- item 1
 - another item
- item 2
- item 3

Numbered lists

1. item 1
2. item 2
3. item 3

The text body

Math expressions can be written in LaTeX syntax

- Inline (enclosed with \$)

```
$f(k) = {n \choose k} p^k (1-p)^{n-k}$ returns  $f(k) = \binom{n}{k} p^k (1 - p)^{n-k}$ 
```

- As a separate block (enclosed with \$\$)

$$f(k) = \binom{n}{k} p^k (1 - p)^{n-k}$$

The text body

You don't need to remember all of this. [Here](#) is a quick reference for the most important things.

Always use spaces around markdown objects so that they can be rendered correctly, e.g.

```
## My section
```

This is an ordered list:

1. item 1
2. item 2

instead of

```
## My section
```

This is an ordered list:

1. item 1
2. item 2

The text body

Let's look at an example together

Can you identify the formatting elements in [this document](#)?

Now you

Task 1: Create some markdown text

Find the task description [here](#)

The R code

The R code

R code can be included in **code chunks** as **inline code**

- R code can be displayed or hidden in output document
 - Inline code is hidden
- Code chunks can contain any type of R code
- R code is (by default) executed and output is included in document
 - Text output
 - Figures
 - Prepare future data analysis
 - ...

The R code

Code chunks starts and ends with 3 backticks

```
```{r}
code goes here
```
```

Example

```
```{r}
mean(1:3)
```
```

looks like this:

```
mean(1:3)

## [1] 2
```

The R code

Inline code starts and ends with 1 backtick

```
## `r `
```

Example

```
The mean of the values 1, 2 and 3 is `r mean(1:3)`
```

looks like this:

The mean of the values 1, 2 and 3 is 2.

The R code

Insert a code chunk

- Insert a code chunk by going to `Code -> Insert chunk`
- Use the keyboard shortcut `Ctrl + Alt + I / Cmd + Option + I`
- Inline chunks have to be typed

The R code

Chunk names

- Code chunks can have names (but they don't have to)
- Names are added inside the code chunk

```
```{r calculate-mean}  
mean(1:3)
```
```

- Easier to debug
- Code chunks appear in document outline

The R code

Run code chunk

- Code chunks are evaluated by `knitr` when rendering the document
- Code chunks can also be run like normal R code
- Run Code chunk by clicking on the green arrow next to the chunk

```
```{r cars}
summary(cars)
```
```

Chunk options

- Code chunks options give you fine control over the behavior of a chunk
- Chunk options are separated by commas

```
```{r calculate-mean, warning=TRUE, message=TRUE}
mean(1:3)
```
```

- Chunk options have a `name` and a `value`
 - Chunk options have default values
- Different document outputs support different chunk options

Chunk options

Some important general options

- `message`: TRUE, FALSE, Show message in output?
- `warning`: TRUE, FALSE, Show warning in output?
- `eval`: TRUE, FALSE, Evaluate the chunk?
- `echo`: TRUE, FALSE, Show source code in output?
- `include`: TRUE, FALSE, Include anything from the chunk in the output?
 - Runs the code but does not show any output or source code

Chunk options

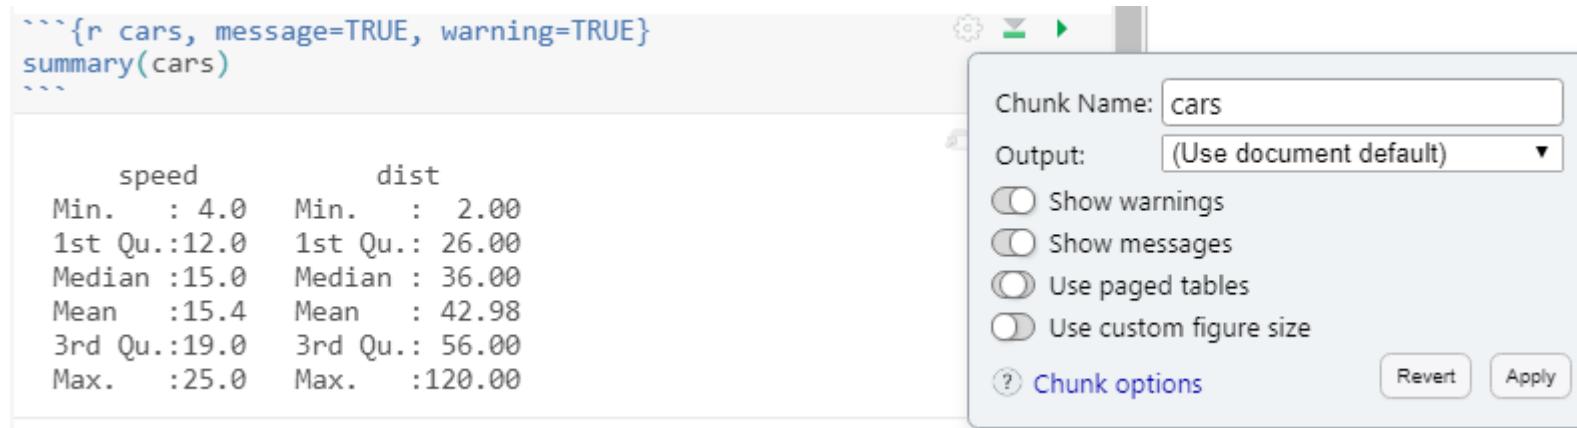
Some important figure related options

- `fig.width` and `fig.height`: Size of graphical device in inches (i.e. size of the plots)
- `out.width` and `out.height`: Scale output of R plots, e.g. to scale images `out.width = "80%"`
- `fig.align`: Plot alignment, one of `"left"`, `"center"`, `"right"`
- `fig.cap`: A figure caption as a string

```
```{r important-figure, fig.cap="This is a nice plot.", fig.width=4, fig.height=4}
plot(1:10, 1:10)
```
```

Chunk options

- See [here](#) for a comprehensive list of all chunk options
- Some chunk options can be set by clicking on the gear icon next to the chunk



A screenshot of the RStudio interface illustrating chunk options. On the left, a code chunk displays the R command `summary(cars)` and its output, which is a summary statistics table for the 'cars' dataset:

```
```{r cars, message=TRUE, warning=TRUE}
summary(cars)
```


	speed	dist
Min. :	4.0	2.00
1st Qu.:	12.0	26.00
Median :	15.0	36.00
Mean :	15.4	42.98
3rd Qu.:	19.0	56.00
Max. :	25.0	120.00


```

On the right, a gear icon is expanded to show the 'Chunk options' dialog box. The 'Chunk Name' field is set to 'cars'. The 'Output' dropdown is set to '(Use document default)'. Below these, four checkboxes are present: 'Show warnings', 'Show messages', 'Use paged tables', and 'Use custom figure size'. At the bottom of the dialog are 'Revert' and 'Apply' buttons.

The setup chunk

- Set default chunk options in the beginning of the document
- Default options can be overwritten in individual chunks
- Default chunk options can be set via `knitr` in an R code chunk:

```
```{r setup, include=FALSE}
knitr::opts_chunk$set(
 echo = TRUE,
 warning = FALSE
)
```

```

Other language engines

You can also run code chunks of other languages:

```
names(knitr::knit_engines$get())
```

```
## [1] "awk"      "bash"     "coffee"    "gawk"     "groovy"    "haskell"   "lein"
## [8] "mysql"    "node"     "octave"    "perl"     "pgsql"     "Rscript"    "ruby"
## [15] "sas"       "scala"    "sed"       "sh"       "stata"     "zsh"        "asis"
## [22] "asy"       "block"    "block2"    "bslib"    "c"         "cat"        "cc"
## [29] "comment"   "css"      "dot"       "embed"    "fortran"   "fortran95" "go"
## [36] "highlight" "js"       "julia"     "python"   "R"         "Rcpp"       "sass"
## [43] "scss"      "sql"      "stan"      "targets"  "tikz"     "verbatim"
```

See [here](#) if you are interested

Metadata with the YAML header

YAML (/ˈjæməl/) header

- YAML is a data-serialization language
- Goes between `---` on top
- Used for
 - Meta data
 - Document output formats and their options
- Is used by Pandoc, `rmarkdown` and `knitr`

YAML (/ˈjæməl/) header

Example

```
---
```

```
title: "My first document"
author: "Selina Baldauf"
date: "3/22/2022"
output:
  html_document:
    toc: true
    toc_float: true
bibliography: references.bib
---
```

- YAML Syntax: Indentation is important
 - Two spaces or a tab for lists
 - Values follow a :
- Translation:
 - TRUE, FALSE, NULL become true, false, null
 - Strings can be quoted but don't have to (if they don't contain special characters)
- Can also contain inline code to be evaluated

YAML (/ˈjæməl/) header

Metadata

```
---
```

```
title: "My first document"
subtitle: "Whatever subtitle makes sense"
author: "Selina Baldauf"
date: "`r Sys.Date()`"
```

```
---
```

- Inline R code can print the current date at knitting time

YAML (/ˈjæməl/) header

Document outputs

- html_document
- pdf_document
- word_document
- beamer_presentation
- powerpoint_presentation
- ...

To use the default version of an output format add

```
output: html_document
```

or

```
output:  
  html_document: default
```

YAML (/ˈjæməl/) header

Document output options

- Every output format has options than can be specified in the YAML header
 - Some options are shared between formats
 - Some options are specific to formats
- Have a look at `?rmarkdown::pdf_document`, `?rmarkdown::html_document`, etc. to see all options and their default values
- Use the [rmarkdown cheat sheet](#) as reference for the most important options

YAML (/ˈjæməl/) header

Document output options

```
output:  
  html_document:  
    toc: true  
    toc_depth: 2  
    number_sections: true  
    fig_caption: true  
    highlight: "espresso"
```

- `toc`: Show table of contents?
- `toc_depth`: Lowest header level for table of contents
- `number_sections`: Should sections be numbered?
- `fig_caption`: Render figure captions?
- `highlight`: Which style to use for syntax highlighting?
 - Some options: "kate", "tango", "zenburn"

YAML (/ˈjæməl/) header

Document parameters

```
params:  
  species: "setosa"
```

Document parameter values can be used in R code (inline or code chunks) with `params$param_name`

```
```{r iris-figure, fig.width=4, fig.height=4}  
to_plot <- subset(iris, Species == params$species)
plot(Sepal.Length, to_plot, data = to_plot)
```
```